

Cloud Computing Tutorial Session 4



Björn Lohrmann

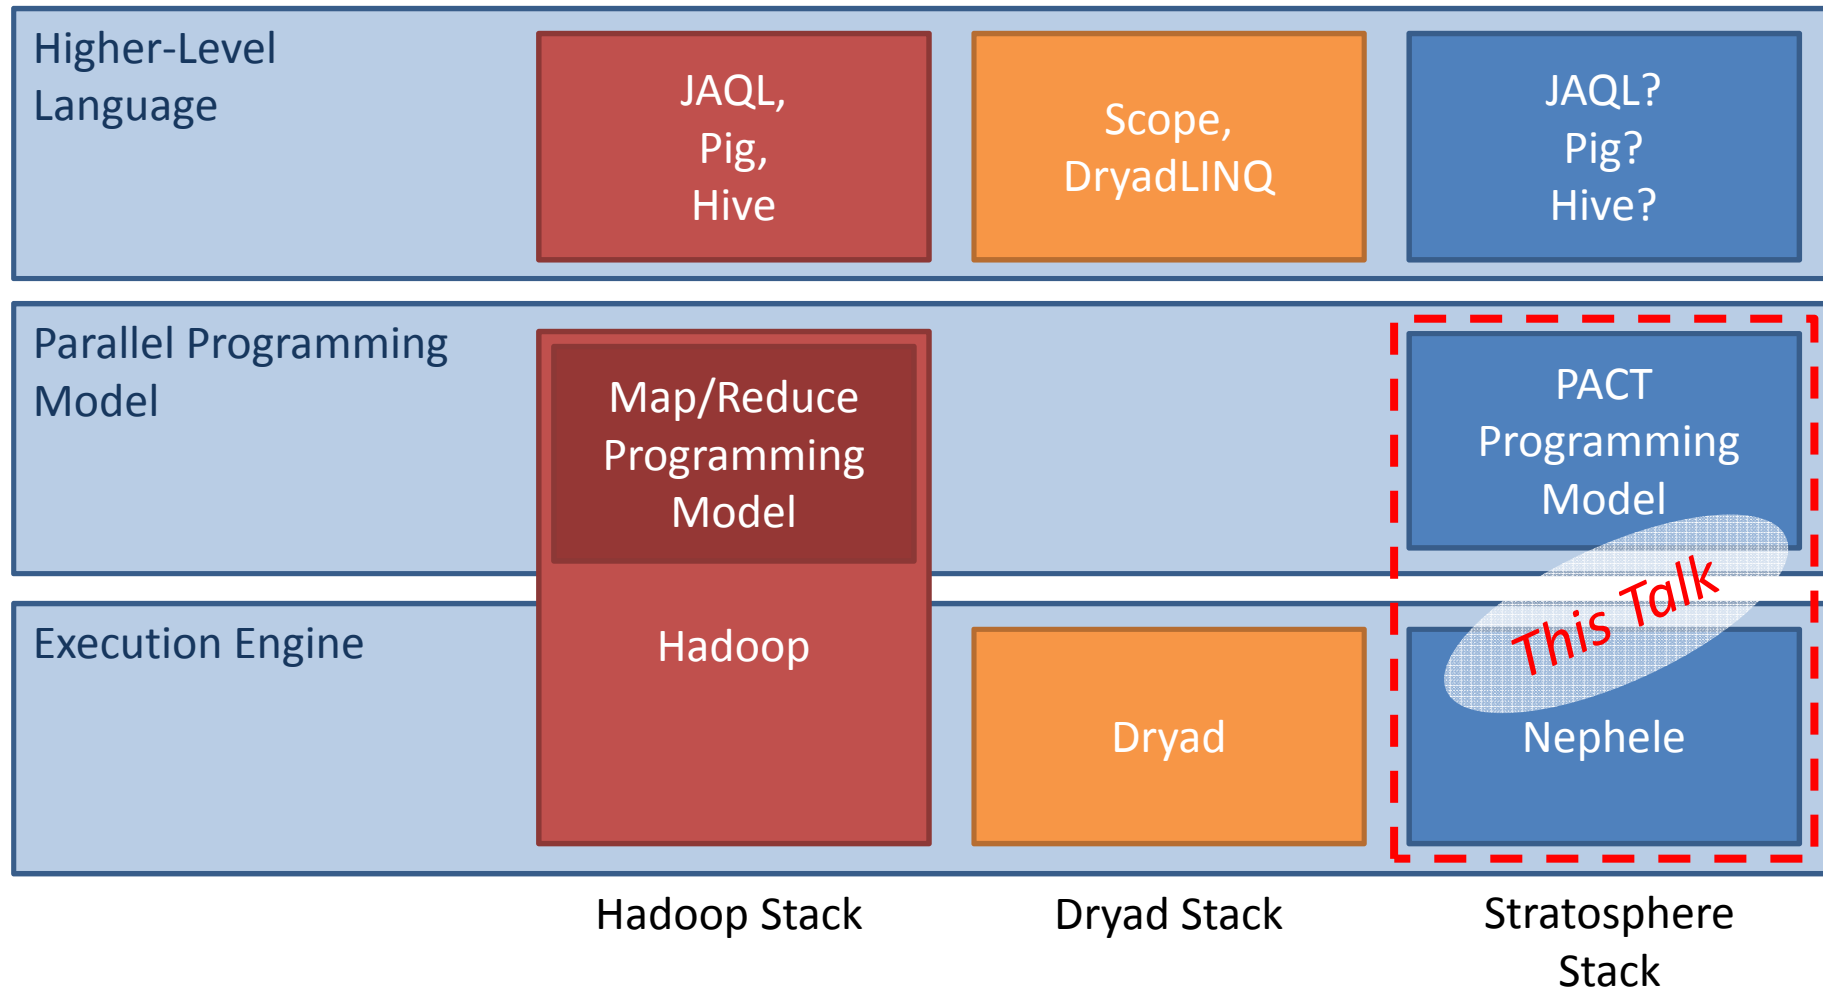
Complex and Distributed IT-Systems

Bjoern.lohrmann@tu-berlin.de

Agenda


- Architecture of the Stratosphere System
 - The PACT Programming Model
 - The Nephele Execution Engine
- Project Assignment #4
- TPC-H benchmark data

Architecture Overview



Stratosphere in a Nutshell

- PACT Programming Model
 - Parallelization Contract (PACT)
 - Declarative definition of data parallelism
 - Centered around second-order functions
 - Generalization of map/reduce
- Nephelē
 - Dryad-style execution engine
 - Evaluates dataflow graphs in parallel
 - Data is read from distributed filesystem
 - Flexible engine for complex jobs
- Stratosphere = Nephelē + PACT
 - Compiles PACT programs to Nephelē dataflow graphs
 - Combines parallelization abstraction and flexible execution
 - Choice of execution strategies gives optimization potential




An Intuition for Parallelization Contracts (PACTs)

- Map and reduce are second-order functions
 - Call first-order functions (user code)
 - Provide first-order functions with subsets of the input data

- Map and reduce are PACTs in our context


- Map

- All pairs are independently processed



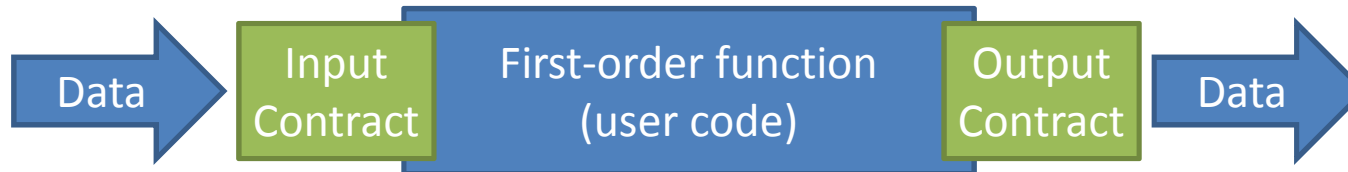
- Reduce

- Pairs with identical key are grouped
- Groups are independently processed



What is a PACT?




- Second-order function that defines properties on the input and output data of its associated first-order function



- Input Contract
 - Generates independently processable subsets of data
 - Generalization of map/reduce
 - Enforced by the system
- Output Contract
 - Describes properties of the output of the first-order function
 - Use is optional but enables certain optimizations
 - Guaranteed by the user
- Key-Value data model

PACTs beyond Map and Reduce

- Cross
 - Multiple inputs
 - Cartesian Product of inputs is built
 - All combinations are processed independently
- Match
 - Multiple inputs
 - All combinations of pairs with identical key over all inputs are built
 - All combinations are processed independently
 - Contract resembles an equi-join on the key
- CoGroup
 - Multiple inputs
 - Pairs with identical key are grouped for each input
 - Groups of all inputs with identical key are processed together



Output Contracts *(examples)*

- Same-Key

- User Function does not alter the key
- For Multi-Input PACTs specify whose input-key remains



- Super-Key

- Key generated by UF is a super-key of the input key
- For Multi-Input PACTs specify from which input the key is a super-key




- Unique-Key

- UF produces unique keys



Nephele Execution Engine



- Executes Nephele schedules
 - compiled from PACT programs
- Design goals
 - Exploit scalability/flexibility of clouds
 - Provide predictable performance
 - Efficient execution on 1000+ nodes
 - Introduce flexible fault tolerance mechanisms
- Inherently designed to run on top of an IaaS Cloud
 - Can exploit on-demand resource allocation
 - Heterogeneity through different types of VMs possible
 - Knows Cloud's pricing model



Nephele Architecture

- Standard master worker pattern
- Workers can be allocated on demand

Workload over time



Project Assignment #4

- **Goal**
 - Compare Hadoop and PACT programming models for a simple query
 - Use Hadoop and Stratosphere frameworks
- **Test data**
 - TPC-H benchmark tools
 - Generates relational data model in ASCII format

- **Relations relevant for assignment**
 - **Customers**
 - Have key, name and other attributes
 - **Orders**
 - Customers create orders
 - Has key, totalprice and other attributes
 - Has foreign key on customer
 - **Lineitems**
 - Orders consist of lineitems
 - Has foreign key on order